

# Offroad Semantic Segmentation - Hackathon Report

**Team Leader:** Karthik Rajarapu

**Team Members:** S. Soheb, Binita, Divya

**Hackathon:** Duality AI - Offroad Autonomy Segmentation Challenge

**Date:** January 31, 2026

## Executive Summary

This report documents our approach to training a semantic segmentation model for offroad terrain analysis using Duality AI's synthetic desert environment dataset. Our best model achieved **55.5% IoU** (57.2% with TTA) using a custom UNet architecture with class-weighted CrossEntropy loss.

Metric	Value
<b>Final IoU Score</b>	<b>57.2%</b> (with TTA)
<b>Model Architecture</b>	Custom UNet (4 encoder, 4 decoder blocks)
<b>Training Time</b>	1.5 hours (T4 GPU)
<b>Experiments Conducted</b>	7

## 1. Methodology

### 1.1 Data Preprocessing

**Challenge:** Raw images needed standardization for efficient training.

**Solution:** Offline preprocessing pipeline:

- Resized all images to 256×256
- Converted masks to indexed format (0-9 class IDs)
- Organized into train/val splits (80/20)

**Result:** 95% reduction in I/O time during training.

### 1.2 Model Architecture

We implemented a **Custom UNet** with:

- **Encoder:** 4 blocks (64→128→256→512 channels)
- **Decoder:** 4 blocks with skip connections
- **Output:** 10-class segmentation map

```
Input (3x256x256)
  ↓
Encoder: Conv→BN→ReLU→MaxPool (x4)
```

```

↓
Decoder: Upsample→Concat→Conv→BN→ReLU (×4)
↓
Output (10×256×256)

```

### 1.3 Training Configuration

Parameter	Value
Loss Function	CrossEntropyLoss (weighted)
Optimizer	Adam (lr=0.001)
Batch Size	8
Epochs	50
Hardware	NVIDIA T4 GPU

### 1.4 Class Weighting Strategy

**Challenge:** Severe class imbalance (e.g., Logs: 0.08% vs Landscape: 38%)

**Solution:** Inverse-frequency class weighting:

Class	Pixel %	Weight Applied
Landscape	38%	0.3×
Trees	3.6%	2.8×
<b>Logs</b>	<b>0.08%</b>	<b>130×</b>

## 2. Results & Performance Metrics

### 2.1 Final IoU Score

Split	IoU
<b>Validation</b>	<b>55.5%</b> (57.2% with TTA)

### 2.2 Experiments Summary

#	Model	Configuration	IoU	Status
1	Custom UNet	CE + Class Weights	<b>55.5%</b> (57.2% TTA)	<input checked="" type="checkbox"/> Best
2	Custom UNet	Dice+CE Combined	40.6%	✗
3	ResNet34-UNet	Pretrained Encoder	45.0%	✗
4	Custom UNet	100 epochs	51.1%	△

#	Model	Configuration	IoU	Status
5	ResNet34-UNet	Frozen Encoder	40.5%	✗
6	Custom UNet	512×512 resolution	20.4%*	⚠️
7	DeepLabV3+	Focal+Dice Loss	N/A	✗ Crashed

\*Only 1 epoch completed before switching strategy

## 2.3 Training Progression

```
Epoch 10: IoU = 0.37
Epoch 25: IoU = 0.48
Epoch 40: IoU = 0.53
Epoch 50: IoU = 0.555 (Base)
With TTA: IoU = 0.572 (Final)
```

## 2.4 Per-Class Performance (Estimated)

Class	Expected IoU	Notes
Landscape	High	Largest class, well-learned
Sky	High	Large, distinctive
Trees	Medium	Good features
Dry Grass	Medium	Well-represented
<b>Logs</b>	Low	Very rare (0.08%)
<b>Rocks</b>	Low	Occlusion issues

## 3. Challenges & Solutions

### Challenge 1: Severe Class Imbalance

**Problem:** Logs class represents only 0.08% of pixels, causing the model to ignore it entirely.

**Solution:** Implemented inverse-frequency class weighting:

```
class_weights = [total_pixels / (num_classes * class_pixels)]
```

**Result:** Rare classes became detectable; overall IoU improved from 0.31 to 0.572.

### Challenge 2: Slow Training (60 min/epoch)

**Problem:** CPU-based training with on-the-fly preprocessing was too slow.

**Solution:**

1. Migrated to GPU (T4)
2. Offline preprocessing (256×256 resizing)
3. Optimized DataLoader (pin\_memory, num\_workers)

**Result:** Training speed improved **50x** (60 min → <2 min per epoch).

---

### Challenge 3: Pretrained Models Underperformed

**Problem:** ResNet34 pretrained on ImageNet achieved lower IoU (45%) than custom UNet (57.2%).

**Analysis:**

- ImageNet features (cats, dogs) don't transfer well to desert terrain
- Domain gap between natural images and synthetic desert data
- ImageNet normalization mismatch

**Lesson:** For domain-specific synthetic data, training from scratch can outperform pretrained models.

---

### Challenge 4: Dice Loss Instability

**Problem:** Dice+CE combined loss led to oscillating training and lower final IoU (40.6%).

**Analysis:**

- Dice loss gradient signals conflict with CE
- Harder to optimize than pure CE
- Required more careful hyperparameter tuning

**Solution:** Reverted to simple CE with class weighting.

---

## 4. Failure Case Analysis

Misclassification Patterns Observed:

Failure Type	Likely Cause	Potential Fix
Small logs missed	Low resolution (256×256)	Train at 512×512
Rock/Ground confusion	Similar texture	Add texture augmentation
Edge imprecision	Downsampling artifacts	Use higher resolution
Shadow misclassification	Color variation	Brightness augmentation

---

## 5. Optimizations Applied

### 5.1 Test-Time Augmentation (TTA)

Implemented TTA for inference-time accuracy boost:

- Horizontal flip
- Vertical flip
- Average predictions from 3 views

**Expected Improvement:** +2% IoU with zero retraining.

## 5.2 Data Pipeline Optimization

Optimization	Impact
Offline preprocessing	95% I/O reduction
GPU migration	50× speedup
Efficient DataLoader	2× throughput

## 6. Conclusion & Future Work

Key Findings:

1. **Simple architectures win:** Custom UNet outperformed pretrained ResNet34
2. **Class weighting is critical:** Enables detection of rare classes
3. **Domain matters:** ImageNet pretraining doesn't help with synthetic desert data
4. **Preprocessing pays off:** Offline data preparation is worth the investment

Future Improvements:

- Higher resolution training (512×512) with mixed precision
- Focal Loss for hard example mining
- Ensemble of multiple models
- CRF post-processing for boundary refinement
- Domain-specific augmentation (desert shadows, dust effects)

## 7. Reproducibility Instructions

Environment Setup:

```
pip install torch torchvision tqdm pillow numpy
```

Training:

```
python train.py
```

Inference:

```
python test.py
```

## Files Included:

- `train.py` - Training script with class weighting
  - `test.py` - Inference script with TTA
  - `checkpoints/best_model.pth` - Trained model weights
  - `README.md` - Quick start guide
- 

**Thank you for reviewing our submission!**

*Team Leader: Karthik Rajarapu*

*Team Members: S. Soheb, Binita, Divya*